

Q1

Q1.1.

```
[1]: import matplotlib.pyplot as plt
import numpy as np


class F:
    const = [10,1,1] #a,b,c

    # calc f(x)
    @staticmethod
    def f(x):
        y = 0
        for i, const in enumerate(F.const):
            y += (x**i)*const
        return y

    # calc gradient of f(x)
    @staticmethod
    def grad_f(x):
        y = 0
        for i in range(1, len(F.const)):
            y += i * F.const[i] * (x ** (i - 1))
        return y
```

```
[2]: f_x = np.linspace(-10, 10, 400)
f_y = F.f(f_x)

plt.plot(f_x, f_y, '-', label='f(x) values')
plt.xlabel('x')
plt.ylabel('y')
plt.title('f(x) = x^2 + x + 10')
plt.legend()
plt.show()
```



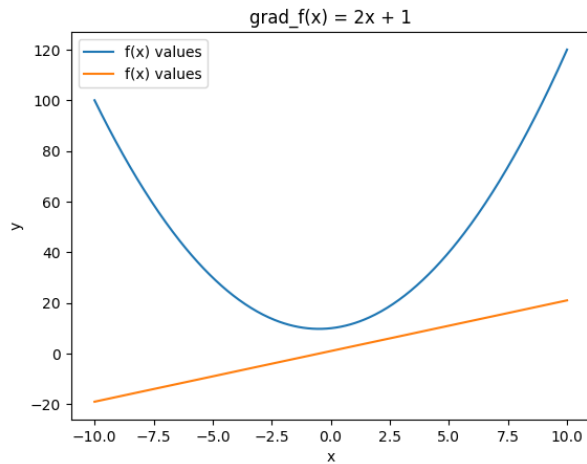
Q1.2.

```
[3]: f_x = np.linspace(-10, 10, 400)
     f_y = F.f(f_x)

     grad_f_x = np.linspace(-10, 10, 400)
     grad_f_y = F.grad_f(grad_f_x)

     plt.plot(f_x, f_y, '-', label='f(x) values')
     plt.plot(grad_f_x, grad_f_y, '-', label='f(x) values')

     plt.xlabel('x')
     plt.ylabel('y')
     plt.title('grad_f(x) = 2x + 1')
     plt.legend()
     plt.show()
```



Q1.3.

$\min(f(x))$: $\text{grad}_f(x) = 2x + 1 = 0 \rightarrow x = -0.5, y = 9.75$

Q1.4.

```
[4]: def GD(init_x, epsilon, n, save_x):
     x = []
     x_t = float('inf') # x_(t)
     x_T = init_x # x_(t+1)
     t = 0 # Iteration counter

     while abs(x_t - x_T) > epsilon:
         x_t = x_T
         x_T -= n * F.grad_f(x_T)
         if save_x: x.append(x_T)
         t += 1
         if t > 1000: # Prevent infinite loops by setting a maximum number of iterations
             print("Maximum iterations reached")
             break

     return x_T, t, x
```

Q1.5.

```
[5]: # Running the gradient descent algorithm
     init_x = 10
     epsilon = 0.00001
     n = 0.1 # Learning rate

     result, t, x = GD(init_x, epsilon, n, False)
     print(f"The minimum is aproxemtry at x = {result}, calc toke t = {t} iterations")

     The minimum is aproxemtry at x = -0.4999607148359886, calc toke t = 56 iterations

     the value we got is not equal to the minimum of f(x) because of the epsilon we defines as 0.00001
```

Q1.6.

```
[6]: # Running the gradient descent algorithm
init_x = 10
epsilon = 0.01
n = 0.4 # Learning rate

result, t, x = GD(init_x, epsilon, n, True)
print(f"The minimum is aproxemtry at x = {result}, calc toke t = {t} iterations")


The minimum is aproxemtry at x = -0.499328, calc toke t = 6 iterations
```

Q1.7.

```
[7]: f_x = np.linspace(-10, 10, 400)
f_y = F.f(f_x)

x = np.array(x)
y = F.f(x)

plt.plot(f_x, f_y, '-', label='f(x) values')
plt.plot(x, y, 'o', label='f(x) values')
plt.xlabel('x')
plt.ylabel('y')
plt.title('f(x) = x^2 + x + 10')
plt.legend()
plt.show()
```



Q2

Q2.1.

target func is: $SVG(w, x, y) = \arg\min_{w \in \mathbb{R}^d, b \in \mathbb{R}} \sum (\max\{0, 1 - y_i(\langle w, x_i \rangle + b)\} + \lambda \|w\|^2)$, we will break it down to a few func:

$f(x) = \arg\min(x)$ - convex func from thyrom we learned in TA

$g(x) = \sum(x)$ - convex func from thyrom we learned in TA

$u(x, y) = \max\{x, y\}$ - convex func from thyrom we learned in TA

$h(w, x) = \langle w, x \rangle$ - convex func from thyrom we learned in TA

$k(w) = \|w\|^2$ - norm \rightarrow convex func from thyrom we learned in TA and lemma from HW file

from here: $SGD(w, x, y) = f(g(u(x, 1 - y_i(h(w, x_i) + b) + \lambda k(w))) \rightarrow$ sum of convex funcs** is convex

** linair comb of convex and scalr multi of convex is convex

Q2.2

$$\text{or } f: \mathbb{R}^d \rightarrow \mathbb{R} \text{ is } f \text{ is } L \text{ Lipschitz}$$

$$\|f(w_1) - f(w_2)\| \leq L \|w_1 - w_2\|$$

$$L = \max_k \|x_k\|$$

$$\text{for } f(w, x, y)$$

$$L_1(w_1, x, y), L_2(w_2, x, y) > 0$$

$$\|L_1 - L_2\| = \|y_i \langle w_2, x_i \rangle - y_i \langle w_1, x_i \rangle\|$$

$$\|y_i (\langle w_2, x_i \rangle - \langle w_1, x_i \rangle)\| = \|y_i (\langle w_2 - w_1, x_i \rangle)\|$$

$$\|\langle w_2 - w_1, x_i \rangle\| \leq \|x_i\| \cdot \|w_2 - w_1\|$$

$$\|x_i\| \cdot \|w_2 - w_1\| \leq \max_k \|x_k\| \cdot \|w_2 - w_1\|$$

$$L \cdot \|w_2 - w_1\|$$

$$\leftarrow \mathcal{L}(w_1, x, y), \mathcal{L}(w_2, x, y) > 0 \quad \text{?}$$

$$\|l_1 - l_2\| = \|y_i \langle w_2, x_i \rangle - y_i \langle w_1, x_i \rangle\| :$$

$$\|y_i (\langle w_2, x_i \rangle - \langle w_1, x_i \rangle)\| = \|(\langle w_2, x_i \rangle - \langle w_1, x_i \rangle)\| \cdot \|y_i\|$$

$$\|\langle w_2 - w_1, x_i \rangle\| \leq \|x_i\| \cdot \|w_2 - w_1\| :$$

$$\|x_i\| \cdot \|w_1 - w_2\| \leq \max_x \|x\| \cdot \|w_1 - w_2\| :$$

$$R \cdot \|w_1 - w_2\|$$

$$\leftarrow \mathcal{L}(w_1, x, y), \mathcal{L}(w_2, x, y) = 0 \quad \text{?}$$

$$\|l_1 - l_2\| : 0 - 0 \leq R \cdot \|w_1 - w_2\|$$

$$\leftarrow \text{if } \mathcal{L}(w_1, x, y) > 0, \mathcal{L}(w_2, x, y) = 0 \quad \text{?}$$

$$\|l_1 - l_2\| : 1 - y_i \langle w_1, x_i \rangle - 0 = y_i \langle w_2, x_i \rangle - y_i \langle w_1, x_i \rangle$$

$$\|l_1 - l_2\| \leq R \|w_1 - w_2\| \quad R > 1, \text{ so } \checkmark$$

$$(*) \quad \mathcal{L}_2(w_2, x, y) = 1 - y_i \langle w_2, x, y \rangle = 0 - y_i \langle w_2, x, y \rangle :$$

Q2.3.

As we have seen in the tutorial:

$$(d/dw)SGD(w, x, y) = \{2\lambda w \text{ if } 1 - y_i \langle w, x_i \rangle + b \leq 0, -y_i x_i + 2\lambda w \text{ if } 1 - y_i \langle w, x_i \rangle + b > 0\}$$

$$(d/db)SGD(w, x, y) = \{0 \text{ if } 1 - y_i \langle w, x_i \rangle + b \leq 0, -y_i \text{ if } 1 - y_i \langle w, x_i \rangle + b > 0\}$$

Q2.4.

```
[8]: def svm_with_sgd(X, y, lam=0, epochs=1000, l_rate=0.01, sgd_type='practical'):
    np.random.seed(2)
    m, d = X.shape
    b = np.random.uniform(0, 1, 1)
    w = np.random.uniform(0, 1, d)

    def SG(w, x, y, b, lam):
        if 1 - y * (w.dot(x) + b) <= 0:
            return (2 * lam * w, 0)
        else:
            return (-y * x + 2 * lam * w, -y)

    if sgd_type == 'practical':
        for j in range(epochs):
            per = np.random.permutation(np.arange(0, m))
            for i in range(m):
                k = per[i]
                subgrad = SG(w, X[k], y[k], b, lam)
                w = w - (l_rate * subgrad[0])
                b = b - (l_rate * subgrad[1])
            return (w, b)
    else:
        ws = [w]
        bs = [b]
        for i in range(epochs * m):
            choice = np.random.randint(0, m)
            subgrad = SG(w, X[choice], y[choice], b, lam)
            w = w - (l_rate * subgrad[0])
            b = b - (l_rate * subgrad[1])
            ws.append(w)
            bs.append(b)
        return (sum(ws) / len(ws), sum(bs) / len(bs))
```

Q2.5.

```
[9]: def calculate_error(w,bias,X,y):
    pred = np.add(w.dot(X.T),bias)
    pred[pred> 0] = 1
    pred[pred<=0] = -1
    return np.sum(pred!=y)/X.shape[0]
```

Q2.6.

```
[10]: from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split


# Load data
X, y = load_iris(return_X_y=True)
X = X[y != 0]
y = y[y != 0]
y[y==2] = -1
X = X[:, 2:4]

# split data
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.3, random_state=0)

[11]: lambdas=[0, 0.05, 0.1, 0.2, 0.5]
train_errors = []
test_errors = []
margins = []
for lam in lambdas:
    model = svm_with_sgd(X_train, y_train, lam)
    train_errors.append(calculate_error(model[0],model[1],X_train,y_train))
    test_errors.append(calculate_error(model[0],model[1],X_val,y_val))
    margins.append(1/np.linalg.norm(model[0]))
```


```
[17]: x = np.arange(len(lambdas))
plt.bar(x,train_errors,width=0.35,label='train', color='b')
plt.bar(np.add(x,0.35),test_errors,width=0.35,label='test', color='y')
plt.xlabel('lambda')
plt.ylabel('error')
plt.title('Errors of the sgd model as dependent on lambdas')
plt.xticks(np.add(x,0.35/2),lambdas)
plt.legend()
```

[17]: <matplotlib.legend.Legend at 0x832f520>



```
[18]: x = np.arange(len(lambdas))
plt.bar(x,margins,width=0.40,label='margins', color='b')
plt.xlabel('lambda')
plt.ylabel('margin')
plt.title('Margins of the sgd model as dependent on lambdas')
plt.xticks(x,lambdas)
plt.legend()
```

[18]: <matplotlib.legend.Legend at 0x83547b0>



as we can see, for $\lambda = 0.05$ we get the smallest test and train error. we don't get the smallest margin, but we get the second smallest, therefore we will choose $\lambda = 0.05$

Q2.7.a.

```
[14]: import matplotlib.pyplot as plt
import numpy as np
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split

def svm_with_sgd(X, y, lam=0, epochs=1000, l_rate=0.01, sgd_type='practical'):
    np.random.seed(2)
    m, d = X.shape
    b = np.random.uniform(0, 1, 1)
    w = np.random.uniform(0, 1, d)

    def SG(w, x, y, b, lam):
        if 1 - y * (w.dot(x) + b) <= 0:
            return (2 * lam * w, 0)
        else:
            return (-y * x + 2 * lam * w, -y)

    if sgd_type == 'practical':
        for j in range(epochs):
            per = np.random.permutation(np.arange(0, m))
            for i in range(m):
                k = per[i]
                subgrad = SG(w, X[k], y[k], b, lam)
                w = w - (l_rate * subgrad[0])
                b = b - (l_rate * subgrad[1])
            return (w, b)
    else:
        ws = [w]
        bs = [b]
        for i in range(epochs * m):
            choice = np.random.randint(0, m)
            subgrad = SG(w, X[choice], y[choice], b, lam)
            w = w - (l_rate * subgrad[0])
            b = b - (l_rate * subgrad[1])
            ws.append(w)
            bs.append(b)
        return (sum(ws) / len(ws), sum(bs) / len(bs))

def calculate_error(w, bias, X, y):
    pred = np.add(w.dot(X.T), bias)
    pred[pred > 0] = 1
    pred[pred <= 0] = -1
    return np.sum(pred != y) / X.shape[0]

# Load data
X, y = load_iris(return_X_y=True)
X = X[y != 0]
y = y[y != 0]
y[y == 2] = -1
X = X[:, 2:4]

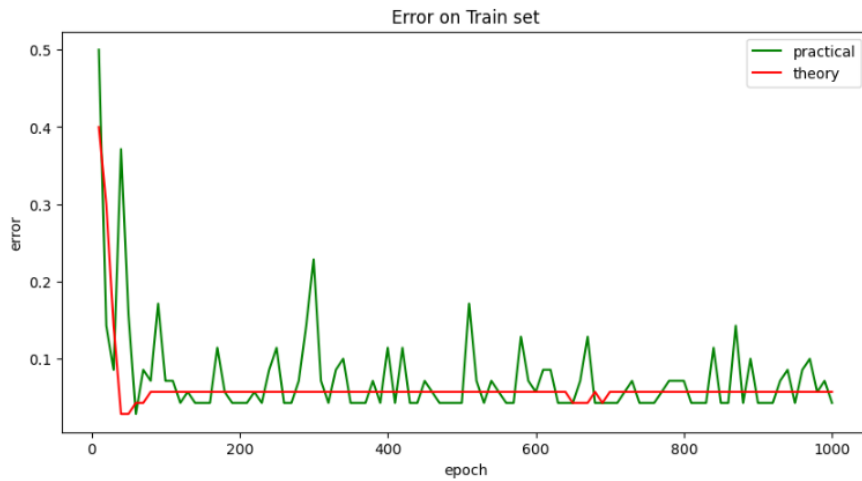
# Split data
X_train, X_val, y_train, y_val = train_test_split(X, y, test_size=0.3, random_state=0)
epochs = np.arange(10, 1001, 10)
train_errors_practical = []
train_errors_theory = []
test_errors_practical = []
test_errors_theory = []

for epoch in epochs:
    modelP = svm_with_sgd(X_train, y_train, lam=0.05, epochs=epoch, sgd_type='practical')
    train_errors_practical.append(calculate_error(modelP[0], modelP[1], X_train, y_train))
    test_errors_practical.append(calculate_error(modelP[0], modelP[1], X_val, y_val))

    modelT = svm_with_sgd(X_train, y_train, lam=0.05, epochs=epoch, sgd_type='theory')
    train_errors_theory.append(calculate_error(modelT[0], modelT[1], X_train, y_train))
    test_errors_theory.append(calculate_error(modelT[0], modelT[1], X_val, y_val))
```




```
[19]: plt.figure(figsize=(10, 5))
plt.plot(epochs, train_errors_practical, label='practical', color='g')
plt.plot(epochs, train_errors_theory, label='theory', color='r')
plt.xlabel('epoch')
plt.ylabel('error')
plt.title('Error on Train set')
plt.legend()
plt.show()
```



Q2.7.b.

```
[20]: plt.figure(figsize=(10, 5))
plt.plot(epochs, test_errors_practical, label='practical', color='g')
plt.plot(epochs, test_errors_theory, label='theory', color='r')
plt.xlabel('epoch')
plt.ylabel('error')
plt.title('Error on Test set')
plt.legend()
plt.show()
```



we can see that in all aproches- theory, practical, train, test we get a highe error fro low epochs that decreases

the theory aproche is more stable then the practical aproche

in the train set we cancle the error in the theory aproche, and in the test set we balnce out around 0.05, likr the λ

3 slide

1 slide

```
import numpy as np
import matplotlib.pyplot as plt
Executed at 2024.07.26 16:15:44 in 12ms
```

```
def cross_validation_errors(X, y, model, folds):
    n = len(X)
    fold_size = n // folds
    train_errors = []
    val_errors = []

    for i in range(folds):
        start = i * fold_size
        end = (i + 1) * fold_size if i < folds - 1 else n # Ensure the last fold includes all remaining data

        X_train = np.concatenate((X[:start], X[end:]))
        y_train = np.concatenate((y[:start], y[end:]))
        X_test = X[start:end]
        y_test = y[start:end]

        # Train the model
        model.fit(X_train, y_train)

        # Predict on training set
        y_train_pred = model.predict(X_train)
        train_error = np.mean(y_train_pred != y_train)
        train_errors.append(train_error)

        # Predict on validation set
        y_test_pred = model.predict(X_test)
        val_error = np.mean(y_test_pred != y_test)
        val_errors.append(val_error)

    average_train_error = np.mean(train_errors)
    average_val_error = np.mean(val_errors)

    return average_train_error, average_val_error
Executed at 2024.07.26 16:15:44 in 28ms
```

```
import sklearn.svm as svm
```

! 2 2.00

```
def svm_results(X_train, y_train, X_test, y_test):
```

```
    results = {}
```

```
    lambda_values = [0.0001, 0.01, 1, 100, 10000]
```

```
    for lambda_val in lambda_values:
```

```
        # Define the SVM model with the given lambda
```

```
        model = svm.SVC(C=(1/lambda_val), kernel='linear') # Example with polynomial kernel, adjust as needed
```

```
        # Calculate training and validation errors using cross-validation
```

```
        avg_train_error, avg_val_error = cross_validation_errors(X_train, y_train, model, folds=5)
```

```
        # Fit the model on the entire training set and compute test error
```

```
        model.fit(X_train, y_train)
```

```
        y_test_pred = model.predict(X_test)
```

```
        test_error = np.mean(y_test_pred != y_test)
```

```
        # Store the results in the dictionary
```

```
        results[f'lambda_{lambda_val}'] = (avg_train_error, avg_val_error, test_error)
```

```
    return results
```

```
from sklearn.datasets import load_iris
iris_data = load_iris()
X, y = iris_data['data'], iris_data['target']
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=7)
```

Executed at 2024.07.26 16:15:44 in 58ms

Get the results

```
results = svm_results(X_train, y_train, X_test, y_test)
```

Prepare data for plotting

```
lambda_values = [0.0001, 0.01, 1, 100, 10000]
train_errors = [results[f'lambda_{lambda_val}'][0] for lambda_val in lambda_values]
val_errors = [results[f'lambda_{lambda_val}'][1] for lambda_val in lambda_values]
test_errors = [results[f'lambda_{lambda_val}'][2] for lambda_val in lambda_values]
```

Plotting

```
x = np.arange(len(lambda_values)) # the label locations
width = 0.2 # the width of the bars
```

```
fig, ax = plt.subplots()
bars1 = ax.bar(x - width, train_errors, width, label='Train Error', color='blue')
bars2 = ax.bar(x, val_errors, width, label='Validation Error', color='orange')
bars3 = ax.bar(x + width, test_errors, width, label='Test Error', color='green')
```

Text for labels, title and x-axis tick labels

```
ax.set_xlabel('Lambda')
ax.set_ylabel('Error')
ax.set_title('Errors for Different Lambda Values')
ax.set_xticks(x)
ax.set_xticklabels(lambda_values)
ax.legend()
```

```
fig.tight_layout()
```

```
plt.show()
```

:3 3.80

שאלה 4

נתון: $g(w) = \max_{i \in [r]} g_i(w)$

עם i , g_i קאונד אפונד \mathbb{R}^d .

$$j \in \arg \max_i g_i(w)$$

צב"ל: $\nabla g_i(w)$ סאב גראדנט g ב- w .

נתון $j \in \arg \max_i g_i(w)$ אופן $g(w) = \max_{i \in [r]} g_i(w) = g_j(w)$ *
 אי-סיוון הסטאנדרט (נתון g אפונד וקאונד)

ובי $u \in \mathbb{R}^d$. מתקיים עבורו:

$$g(u) \geq g(w) + \langle u - w, \nabla g(w) \rangle = g(w) + \langle u - w, \nabla (\max_{i \in [r]} g_i(w)) \rangle$$

האזרה $\nabla g(w)$

$$= g(w) + \langle u - w, \nabla g_j(w) \rangle$$

$$* \leq \langle \nabla g_j(w), u - w \rangle$$

ע"פ האזרה